

Introduction:

This is a note to myself. But I would really appreciate feedback, because this is something that really makes me uncomfortable. Like this presents the Binary Index Tree scheme that I can prove runtimes for but feel shaky as to well how to arrive at it. It's a topic I've seen on LeetCode and covered in fear.

Problem Statement:

2 approaches prefix sum and vanilla list each with tradeoffs. Reconciliation, say we have equal, 50-50 split of Range and Update operations.

Premise:

Let's work in a "universe" of a list of size 15, 1-indexed. All numbers, then, will be from 1-15 and can be expressed using 4 bits from 0b0001 to 0b1111. And Log will mean Log base 2. So given I'm using 4 bits, I have the bound: $\text{Log}(X) \leq 4$. And let the list I'll work with actually be the list : $[1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16]$. This way rather than saying the Ith element of the list I can just say the number I, since the Ith element of the special list IS I.

Preliminary Thoughts

This is just some informal commentary. TODO, really think this out, REQUEST FOR FEEDBACK!

List of size N say I have N buckets. I could give these buckets an ordering and index them as buckets 1,2,...N. Bucket I could definitely have I as a contributor in that bucket and the number I would not show up in any buckets before the Ith bucket (ordering at play)

OK so 2 angles. 1) For computing Range(X), which buckets do we use to compute this? 2) For a given number I, which buckets do we place it in? And for both we want to involve $O(\text{Log}(N))$ buckets, so that the Range and Update operations will be $O(\text{Log}(N))$.

Range LogN

Consider 1) the buckets themselves. I want to introduce 2 concepts: prefixes and branching and then hint (well, blatantly, heavy-handedly) at the BIT scheme satisfying these.

Say we partition X into buckets $\{1-A\} \mid \{A+1 - B\} \mid \{B+1 - X\}$. To compute Range(X) we sum the values of these 3 buckets, where bucket(I - J) holds the sum of $[I + (I+1) + (I+2) + \dots + J]$. For a given X, we want $O(\text{Log}(N))$ buckets

Compare this to the vanilla list approach to compute Range(X): both use prefix sums. In this $\text{Log}(N)$ approach, to get Range(X) I do: bucket $\{B+1 - X\} + \text{Range}(B)$ where, again, bucket $\{B+1 - X\}$ stores the value of the sum of $[B+1, B+2, \dots, X]$. And Range(B) in turn I compute as bucket $\{A+1 - B\} + \text{Range}(A)$ and Range(A) in turn is bucket $\{1-A\}$. So there is a prefix nature of these Range(I) subproblems: Range(A) being a prefix to Range(B) being a prefix to Range(X). I view bucket $\{A+1 - B\}$ as a "child" of bucket $\{1-A\}$ that extends the prefix that is Range(A) to the prefix that is Range(B). And lastly, bucket $\{B+1 - X\}$ as a "child" of bucket $\{A+1 - B\}$ that in turn extends the prefix that is Range(B) to Range(X). So these buckets build upon each other. Now consider the vanilla list $O(N)$ approach to computing Range(X): $\text{Range}(X) = \text{bucket}\{X\} + \text{Range}(X-1)$ where bucket $\{I\}$ is the singleton bucket $\{I - I\}$, that is it simply stores the value I. Now $\text{Range}(X-1)$ in turn I compute as bucket $\{X-1\} + \text{Range}(X-2)$ and so I have $X + (X-1) + \text{Range}(X-2)$ and keep unrolling the Ranges till I get $1 + 2 + \dots + X$. This has the same prefix nature: Range(X-2) is the prefix extending to X-2 and then the singleton bucket $\{X-1\}$ extends that to get Range(X-1) and finally bucket $\{X\}$ extends Range(X-1) to get Range(X). Bucket $\{X\}$ is child of Bucket $\{X-1\}$ is child of Bucket $\{X-2\}$ and so forth.

So both this vanilla list and this $\text{Log}(N)$ approach essentially operate under the same prefix extending mechanism. Main difference seeming to be vanilla list is $O(N)$ and uses X terms/buckets to compute $\text{Range}(X)$ whereas a proposed $\text{Log}(N)$ approach uses only $\text{Log}(N)$ buckets. I will now contemplate potential implications of a $\text{Log}(N)$ scheme and how it would likely be different from the vanilla case: namely, the key idea of branching as opposed to the linear, non-branching approach that is the vanilla list. Upon careful examination of this scheme, I believe it is natural that “branching” happen (TODO think about this, REQUEST FOR FEEDBACK) First point is that under the $\text{Log}(N)$ approach, not all of the buckets can be singleton buckets like the vanilla list case. Because the $\text{Log}(N)$ buckets need to cover all X numbers, $1-X$. Let's say the second bucket $\{A+1 - B\}$ is non-singleton and contains some $J \neq B$, that is, $A+1 \leq J < B$. Important: Consider $\text{Range}(J)$. I believe it is “natural” to reuse the first bucket $\{1-A\}$ to get the prefix of $\text{Range}(A)$ and then take some other path of buckets to get to $\text{Range}(J)$. Note that we can't use the second bucket of $\{A+1 - B\}$ because that would overshoot J . So the idea is that if a child of a bucket B is non-empty, it cannot be an only child and bucket B branches to more than one child where, again, all the children get to reuse the contents of bucket B which in turn reuses of its parents and so forth creating a happy prefix chain (of hopefully length $\leq \text{Log}(N)$).

Actually, just as I write this, I realize this suggests why a binary recursive scheme may be natural. This may be getting a bit ahead of myself, and next paragraph I will give a tiny preview of the BIT scheme, but these prefix and branching notions gave me an idea. TODO review this as it's stream of consciousness, evaluate content/correctness and clarity/notation/try being consistent. Say we follow some Prefix of buckets in a chain that covers $\text{Range}(A)$ and then we reach the bucket $\{A+1 - B\}$ that is non-singleton. So it's parent bucket, P , ends in A and is the last node on the prefix bucket chain that, together, covers $\text{Range}(A)$. From last paragraph discussion on branching to cover all elements, I believe this forces branching. Taking this bucket will jump to B , but intermediate ranges $[A+1 - J]$ for $A+1 \leq J < B$ must be accounted for. Say there are K such J 's ($K = B - 1 - A$) Then say we have some strategy, S , to build a subtree that is also a child of the parent bucket P (so this entire subtree is a sibling to the subtree rooted at bucket $\{A+1 - B\}$) And its job is to cover all these K range queries: $\text{Range}\{A+1\}, \text{Range}\{A+2\}, \dots, \text{Range}\{B-1\}$. Now consider the descendants of the bucket $\{A+1 - B\}$, or the subtree rooted at that bucket. This bucket itself answers $\text{Range}(B)$ when extending its prefix chain and its descendants would answer range queries $\text{Range}(J)$ for $J > B$ for some, but not necessarily all, J . A thought experiment is what if this subtree rooted at the bucket $\{A+1 - B\}$ would answer K more range queries: $\text{Range}\{B+1\}, \text{Range}\{B+2\}, \dots, \text{Range}\{B+K\}$. If I assume the S is an “optimal” strategy for creating a subtree to answer K range queries, it feels natural to me to reuse this same S to create the subtree rooted at $\{A+1 - B\}$. Essentially, this suggests that a recursive approach at least makes sense to me. And a binary one at that because the parent bucket, P , has 2 children: a left subtree (say with a dummy root) that handles K Range queries and a right subtree that is rooted at bucket $\{A+1 - B\}$ which handles $K+1$ Range queries (K for the descendants and $+1$ for $\text{Range}(B)$ that bucket $\{A+1 - B\}$ the root handles).

So for the BIT scheme, we have a bucket $\{1-8\}$ so $\text{Range}(J)$ for $1 \leq J \leq 7$ have some strategy and this same strategy applies for $\text{Range}(K)$ for $9 \leq K \leq 15$. Basically, recursive nature of this arrangement. I'm definitely getting ahead of myself, but I can almost declare completion here itself. $\text{Range}(X)$ is $\text{Log}(N)$ because 3 cases: simplest is $X = 8$ so $\text{Range}(8) = \text{bucket}\{1-8\}$ and done, $1 \leq X < 8$ case recurse the left subtree so instead of universe $[1,15]$ the search is now in universe $[1,7]$ and we have eliminated half the buckets so given there are N buckets total, this will be $O(\text{Log}(N))$ if the recursion keeps eliminating half the buckets. Which it does because the last case is $8 < X \leq 15$ so universe has halved to $[9, 15]$ and $\text{Range}(X)$ is computed in the same manner as would $\text{Range}(X-8)$ except there is an offset of bucket $\{1-8\}$. So it's recursion on the right subtree for $\text{Range}(X)$ is

mechanistically the same as recursion for $\text{Range}(X-8)$ on the left subtree except I have to add bucket{1-8} to the answer.

I'm done with 1) but I just want to throw in a preview of the BIT scheme that respects a potential $\text{Log}(N)$ scheme that uses $\text{Log}(N)$ buckets to partition X . We want to use $\text{Log}(N)$ buckets. Presumably larger X s will require more buckets. Largest number is 15, so hypothetically being "inspired" by binary representation of 15 as 0b1111 and as a preview of the BIT scheme, I'll use the buckets {1-8} | {9-12} | {13-14} | {15} to compute $\text{Range}(15)$ which is exactly what BIT does. Where each 1 in the binary representation has a corresponding bucket with size commensurate with the place of that 1. So take the number 11 or 0b1011. $\text{Range}(11)$ will use buckets {1-8} + {9-10} + {11}. Since these are just preliminary thoughts, the take away is that whatever scheme I use, I need at most $\text{Log}(N)$ buckets which this binary scheme complies with. Why? Well $X \leq N$ has $\log(X)$ bits and at most all of them will be "1"s and in this scheme, the number of buckets $\text{Range}(X)$ uses are the number of 1 bits in the binary representation of X . And note that both $\text{Range}(15)$ and $\text{Range}(11)$ use the same {1-8} bucket and extend it (bucket {1-8} is a parent to bucket {9-12} and {9-10} in the BIT scheme), so this conveniently agrees with aforementioned concepts of reusing prefixes and branching.

So for the BIT scheme, I want to tie this back to recursive intuition, using an "optimal" strategy, S , I presented 2 paragraphs ago to show Range runs in $\text{Log}(N)$. I have a bucket {1-8} so $\text{Range}(J)$ for $1 \leq J \leq 7$ use strategy S and this same S applies for $\text{Range}(K)$ for $9 \leq K \leq 15$. $\text{Range}(X)$ is $\text{Log}(N)$ because 3 cases: simplest is $X = 8$ so $\text{Range}(8) = \text{bucket}\{1-8\}$ and done, $1 \leq X < 8$ case recurse the left subtree so instead of universe $[1,15]$ the search is now in universe $[1,7]$ and we have eliminated half (see next section where I formally develop BIT scheme, again this is a preview) the buckets so given there are N buckets total, this will be $O(\text{Log}(N))$ if the recursion keeps eliminating half the buckets. Which it does because the last case is $8 < X \leq 15$ so universe has halved to $[9, 15]$ and $\text{Range}(X)$ is computed in the same manner as would $\text{Range}(X-8)$ except there is an offset of bucket{1-8}. So it's recursion on the right subtree for $\text{Range}(X)$ is mechanistically the same as recursion for $\text{Range}(X-8)$ on the left subtree except I have to add bucket{1-8} to the answer.

Update LogN

Whew, that finished $O(\text{Log}(N))$ Range . Now consider 2): which buckets we place the index/value I in for another $O(\text{Log}(N))$ operation, Update . So every bucket I place I in, I will have to adjust it's value if I were to execute an $\text{Update}(I)$ value. (lol "I" as in me or "I" as in the variable name I) So at most I can place I in 4 buckets as $\log(16) = 4 > \log(15)$ so the $\text{Update}(I)$ operation will have to touch at most 4 buckets.

Contrast this placing I in $\text{Log}(N)$ buckets approach with the naive prefix sum approach in where instead of $\text{Log}(N)$ buckets, I is placed in $O(N)$ buckets: bucket I , bucket $I+1$, bucket $I+2$, ..., bucket N .

Say I falls in 3 buckets, bucket I , bucket $I+A$, and bucket $I+B$ where bucket K means the K th bucket when the buckets are ordered. Note that I always falls as the highest contributor to the I th bucket. That is $\text{Bucket } I = \text{Bucket}\{J - I\}$ where $J \leq I$. Now when we compute $\text{Range}(X)$ for $X \geq I$, we cannot double count I . So whatever chain of $(\text{Log}(N))$ buckets $\text{Range}(X)$ uses it must use exactly 1 of these buckets. Moreover, from the prefix nature of the buckets (child buckets extending the prefix sum of the chain ending at their parents) none of these 3 buckets can be descendants of each other. Because that would mean double counting I . Basically, avoiding double counting mandates some constraints.

Just like past Range section where I gave a preview of the Range computation for the actual BIT scheme where I used the extreme, highest, number 15, I'll do the same for the Update operation just like how the BIT scheme does it. But this time, the number that appears in the maximum amount of buckets is the extreme, smallest, number that is 1. Which is intuitive because smaller numbers appear in Range sums of all larger numbers so the smallest number will be used the most for all

other numbers. The number 1 appears in bucket{1}, bucket{1-2}, bucket{1-4}, bucket{1-8}. So again, there's this power of 2 deal going on where 1 appears in 4 buckets staying within $\log N$ confines. And when I do Range(X) pursuant to the BIT scheme, I must either get the 1 from bucket{1-8} and potentially continue down that bucket path for all $X \geq 8$ or get the 1 from the remaining buckets. This means $X < 8$ so it's basically binary search, next decision node is if $X \geq 4$ (but < 8) use bucket{1-4}, else $X < 4$ and so forth.

A Binary Tree Scheme

This is the section that essentially should spell everything out exactly, at least for the 4 bit, $N=2^4=16$ case and hopefully for all powers of 2 N . The actual BIT, I maintain, can be viewed as a "compressed" version of this tree where we only focus on non-empty buckets. Consider this guiding picture: start TODO TODO paste this handrawn picture maybe instead of flat files, organize each note into it's own directory to hold the .tpy, .pdf, and image assets together. Maybe also gives room for other file types like animations, TODO/plan files, etc. end TODO